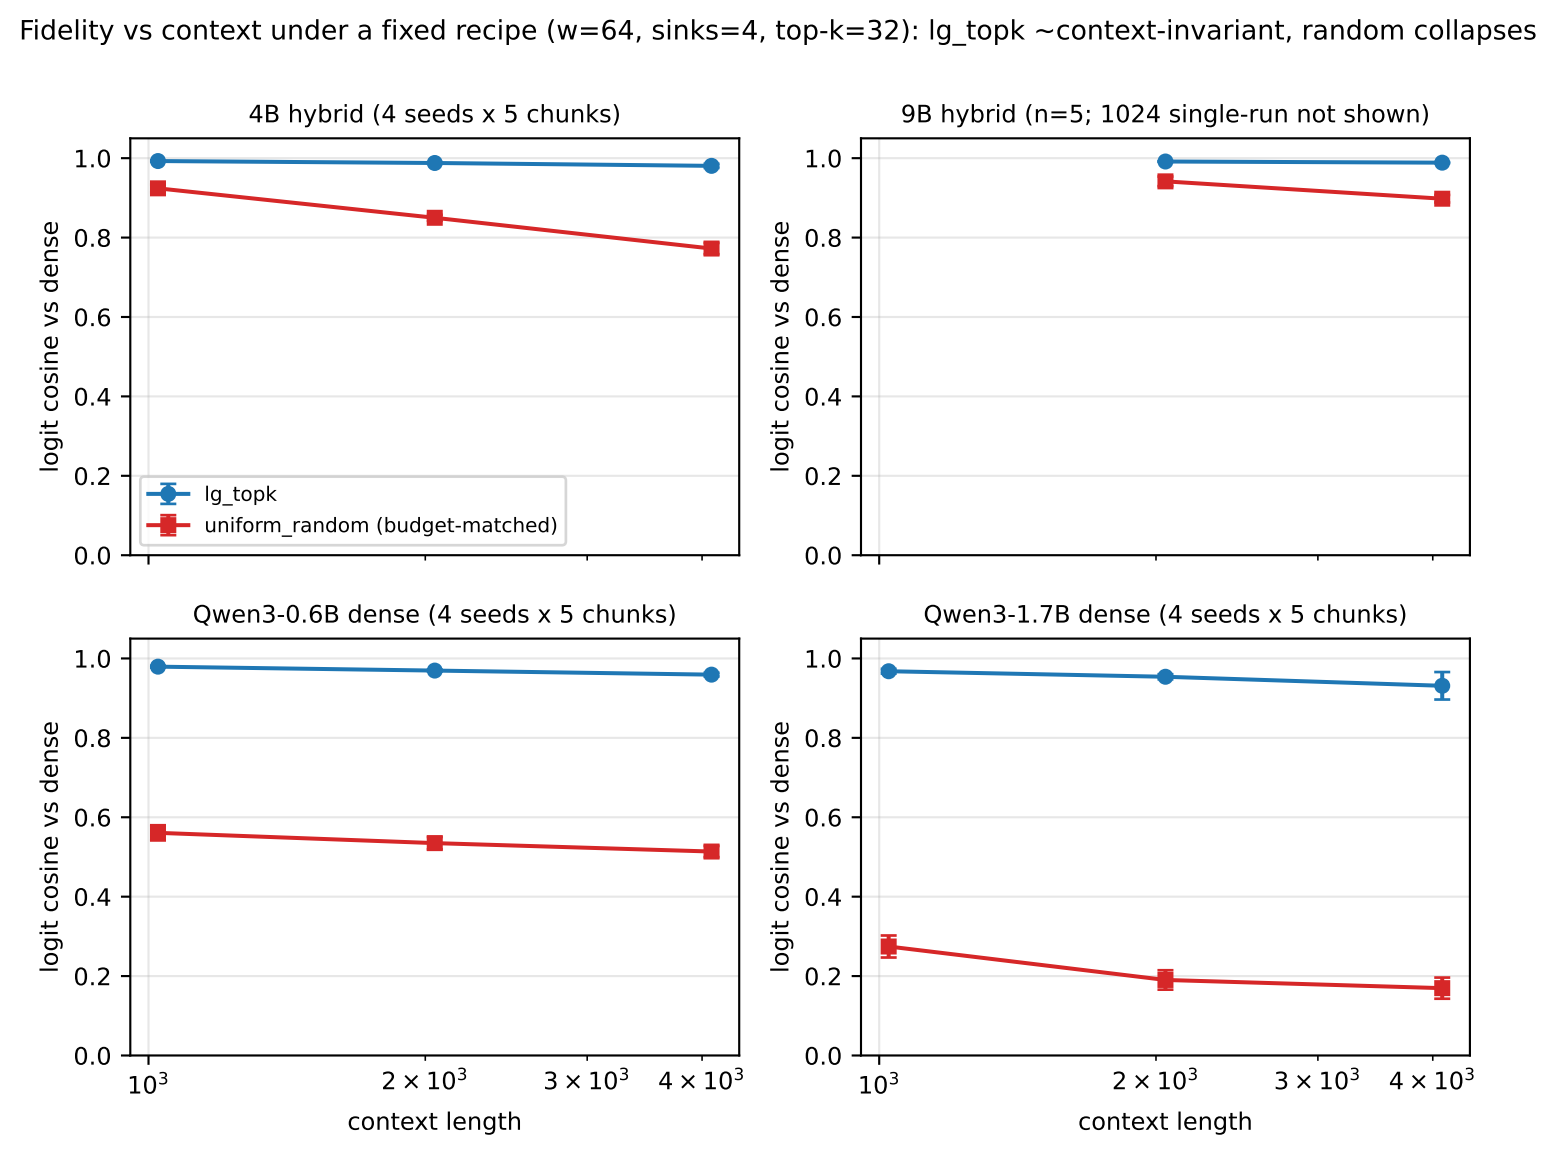
Fidelity vs context under a fixed recipe (w=64, sinks=4, top-k=32): lg_topk ~context-invariant, random collapses
4B hybrid (4 seeds x 5 chunks)
9B hybrid (n=5; 1024 single-run not shown)
Qwen3-0.6B dense (4 seeds x 5 chunks)
Qwen3-1.7B dense (4 seeds x 5 chunks)
logit cosine vs dense
context length
lg_topk
uniform_random (budget-matched)
1.0
0.8
0.6
0.4
0.2
0.0
10^3
2×10^3
3×10^3
4×10^3